

Multicalibration for Unbiased Model-Based Prevalence Estimation

Thomas Leeper¹

Fridolin Linder¹
Lorenzo Perini¹

Daniel Haimovich¹
Milan Vojnovic^{1,2}

Niek Tax¹

¹Meta, ²The London School of Economics and Political Science

Significance

Large language models are increasingly used as measurement devices to estimate prevalence in populations. A critical but overlooked problem arises when the target population differs from the validation population: standard methods produce biased prevalence estimates, even when the model achieves high classification accuracy. We show that multicalibration, requiring a device to be accurate conditional on input features rather than just on average, is sufficient for unbiased prevalence estimation under covariate shift. Our theoretical and empirical results imply that the rapidly growing body of LLM-based measurement research is vulnerable to systematic bias that can be mitigated by enforcing multicalibration.

Abstract

Estimating the prevalence of a category in a population using imperfect measurement devices (diagnostic tests, classifiers, or large language models) is fundamental to science, public health, and platform governance. Standard approaches correct for known device error rates but assume these rates remain stable across populations. We show this assumption fails under covariate shift and that multicalibration, which enforces calibration conditional on the input features rather than just on average, is sufficient for unbiased prevalence estimation under such shift. Standard calibration and quantification methods fail to provide this guarantee. Our work connects recent theoretical work on fairness to a longstanding measurement problem spanning nearly all academic disciplines. A simulation confirms that standard methods exhibit bias growing with shift magnitude, while a multicalibrated estimator maintains near-zero bias. Two empirical applications—estimating employment prevalence across U.S. states using the American Community Survey, and classifying political texts across four countries using an LLM—demonstrate that multicalibration substantially reduces bias in practice, while highlighting that calibration data should cover the key feature dimensions along which target populations may differ.

Introduction

Large language models (LLMs) are increasingly used as measurement devices for estimating the prevalence of a category in a population. Researchers now routinely deploy LLMs as zero-shot classifiers to estimate the frequency of phenomena that previously required expensive manual annotation: examples include, coding democracy indicators across countries (Weidmann et al. 2025), classifying protest events in news corpora (Overös et al. 2024), estimating party policy positions from manifestos across dozens of countries and languages (Benoit et al. 2026), categorizing open-ended survey responses at near-human accuracy (Mellon et al. 2024; Gilardi et al. 2023), extracting diagnostic attributes from pathology reports (Sushil et

al. 2024), identifying goals-of-care discussions in clinical notes (Sibley et al. 2025), annotating art forms in auction records (Tojima and Yoshida 2025), and converting qualitative text into quantitative variables across multiple languages (Karjus et al. 2025). LLMs are also being deployed for content moderation¹ and as LLM “judges” to assess the quality and safety of AI systems (Zheng et al. 2023; Yuan et al. 2024).

These applications are typically validated by reporting discriminative performance: accuracy, AUC, or agreement with human annotations. Researchers then use or advocate for using validated LLMs as measurement devices across new populations, time periods, or subgroups. But strong discriminative performance does not guarantee accurate prevalence estimates. When the composition of the target population differs from the validation population, a model that separates positives from negatives well can still produce prevalence estimates that are substantially biased, because discrimination is invariant to the calibration errors that drive prevalence bias. This problem is invisible to standard validation practice and, as we show, can produce large bias even with strong classification performance.

Neither the confidence elicitation literature nor the quantification² literature solves this problem. Confidence elicitation methods — verbalized confidence (Tian et al. 2023), token log-probabilities (Kadavath et al. 2022), consistency sampling (Wang et al. 2023) — and post-hoc calibration (Oliveira et al. 2025) can improve the quality of an LLM’s scores, but they target *global* calibration: calibration on the population where they are evaluated. Empirical evidence confirms that such calibration does not transfer: it degrades 2–3× under language shift even when accuracy is preserved (Yang et al. 2023), and varies widely across tasks and domains (Ren et al. 2025). Quantification methods such as “Classify & Count”, “Adjusted Count”, and the Saerens-Latinne-Decaestecker EM algorithm (see González et al. (2017) for a detailed survey) attempt to correct for classifier errors directly, but rely on the assumption that error properties remain static across populations. Importance-weighted methods handle covariate shift in principle but require re-estimating density ratios for each new target population. None of these approaches provides a measurement device that can be validated once and then reliably applied to new populations.

We show that multicalibration — calibration conditional on the input features, not just on average — fills this gap. Building on Kim et al. (2022)’s “Universal Adaptability” result, we show that a multicalibrated device requires no target-specific estimation: it is calibrated once on source data and produces unbiased prevalence estimates on any target population whose features lie within the calibration support. This property is critical for LLM-based measurement, where a single device is typically applied to many populations without the opportunity to re-calibrate for each one. Importantly, multicalibration can be applied to any LLM output — from discrete Yes/No classifications (the standard practice in applied research) to continuous confidence scores — because it operates on the scores’ relationship to outcomes conditional on features, not on the scores’ intrinsic quality. Our results apply more generally to any model-based measurement device, but the LLM setting is where the practical need is most acute and the calibration problem most overlooked.

We review existing approaches to the quantification problem, connect the quantification task to the broader theory of domain adaptation via multicalibration, and apply those insights to two empirical applications: a controlled demonstration using a standard ML classifier to estimate employment prevalence under age distribution shift (American Community Survey), followed by the main application of classifying political texts across four countries using an LLM as a zero-shot measurement device.

¹<https://openai.com/index/using-gpt-4-for-content-moderation/>

²The task of estimating prevalence from imperfect classifiers is known as *quantification* in the machine learning literature; we use this term interchangeably with *prevalence estimation*.

Results

Standard calibration methods fail under covariate shift

Consider a binary outcome $Y \in \{0, 1\}$, features X , and a device $h(X) \in [0, 1]$ producing probabilistic predictions. The goal is to estimate population prevalence $\pi = P(Y = 1)$ in a target population using only unlabeled target data $\{X_i\}$ and the device h . We focus on *covariate shift* (Storkey 2009), where $P(X)$ changes across populations but $P(Y | X)$ remains stable. This is the natural assumption when features causally influence outcomes ($X \rightarrow Y$), as in many measurement applications. See Limitations for discussion of label shift ($Y \rightarrow X$) and concept drift.

A device is *globally calibrated* if $\mathbb{E}[Y | h(X) = p] = p$ for all prediction values p . Under global calibration, $\mathbb{E}[h(X)] = \mathbb{E}[Y] = \pi$, so the sample mean of predictions is an unbiased prevalence estimate. However, global calibration is a property of a *specific* population. A device calibrated on one population need not be calibrated on another, even under covariate shift with stable $P(Y | X)$.

The failure mechanism is as follows. Suppose the population consists of subgroups G with weights w_G . Within each subgroup, the device may be biased: let $\epsilon_G = \mathbb{E}[Y - h(X) | X \in G]$ denote the mean prediction error within group G . Global calibration requires only that these errors cancel on average: $\sum_G w_G \epsilon_G = 0$. Under covariate shift, the group weights change to w_G^* , and the bias in the prevalence estimate becomes $\sum_G w_G^* \epsilon_G$, which is generally nonzero unless $\epsilon_G = 0$ for all G . Crucially, a device can be perfectly globally calibrated, with errors that precisely cancel in the training population, while having nonzero mean prediction error within every individual subgroup.

Standard quantification methods are each vulnerable to this failure mechanism (González et al. 2017). Classify & Count, Rogan-Gladen adjustment (Rogan and Gladen 1978), and Probabilistic Adjusted Classify & Count (PACC) all rely on error rates or conditional score means estimated from calibration data; these are weighted averages over subgroups and shift with population composition. The SLD/EMQ algorithm (Saerens et al. 2002) and more recent distribution-matching methods such as DyS and HDy (Maletzke et al. 2020) assume label shift ($P(X | Y)$ stable) and fail under covariate shift. Even global calibration via isotonic regression, recommended by Wu and Resnick (2024) for covariate shift settings, produces biased prevalence estimates because global calibration does not guarantee feature-conditional accuracy.

Multicalibration guarantees unbiased prevalence estimation

The analysis above shows that all standard methods produce biased prevalence estimates under covariate shift, with bias $\sum_G w_G^* \epsilon_G$ that is generally nonzero unless every $\epsilon_G = 0$.

Multicalibration (Hébert-Johnson et al. 2018) formalizes this requirement. A predictor $f(X)$ is *multicalibrated* with respect to a collection of subgroups $\mathcal{G}$ if $\mathbb{E}[Y | f(X) = v, X \in G] = v$ for every $G \in \mathcal{G}$ and prediction value v . When $\mathcal{G}$ is rich enough to capture all relevant structure in X , this is equivalent to requiring calibration conditional on the full feature space:

$$\mathbb{E}[Y | f(X) = v, X = x] = v$$

for all x and v with sufficient probability mass. This is the framing of Kim et al. (2022)’s “Universal Adaptability”: a predictor calibrated conditional on X yields correct expected values under any reweighting of $P(X)$, without requiring knowledge of the shift.

The connection to prevalence estimation is direct. If f is calibrated conditional on X , then $\mathbb{E}[f(X) | X = x] = \mathbb{E}[Y | X = x]$ for all x . Under covariate shift, where only $P(X)$ changes while $P(Y | X)$ remains

stable, the law of iterated expectations gives (where $\mathbb{E}^*$ and P^* denote expectations and probabilities under the target distribution, and $\pi^* = P^*(Y=1)$ is the target prevalence):

$$\mathbb{E}^*[f(X)] = \int \mathbb{E}[f(X) \mid X = x] dP^*(x) = \int \mathbb{E}[Y \mid X = x] dP^*(x) = \mathbb{E}^*[Y] = \pi^*.$$

Because predictions are correct at each point in the feature space, the average of predictions tracks the true prevalence under *any* change in $P(X)$.

The condition strictly necessary for this robustness is *multi-accuracy*: $\mathbb{E}[f(X) - Y \mid X \in G] = 0$ for all groups G that the shift can reweight. Multicalibration, which additionally requires $\mathbb{E}[Y \mid f(X) = v, X \in G] = v$, is strictly stronger and implies multi-accuracy. Multicalibration is the preferable target for three reasons. First, practical post-hoc algorithms like MCGrad (Tax et al. 2026)³ naturally produce multicalibrated predictors (Hébert-Johnson et al. 2018), so the stronger guarantee comes at no additional cost. Second, multicalibration covers a broader set of use cases: when the population shift is mediated by the device’s scores (e.g., score-based sampling for label collection, or threshold-based filtering), when prevalence is estimated within score strata, or when scores are used as inputs to downstream regressions, multi-accuracy no longer suffices and full calibration conditional on both features and score level is required. Third, multicalibration provides robustness against misspecification of which features drive the shift, since a predictor multicalibrated with respect to a rich class of subgroups is automatically multi-accurate for any sub-partition. Both conditions require that the calibration features capture the dimensions along which the population shift occurs (an ignorability assumption).

Simulation: standard methods fail, multicalibration succeeds

We illustrate the theoretical results above with a simulation. We generate data with a binary covariate $X \in \{0, 1\}$ and binary outcome Y , where $P(Y = 1 \mid X = 1) = 0.85$ and $P(Y = 1 \mid X = 0) = 0.15$. A classifier produces systematically biased predictions: 10% underestimation when $X = 0$ and 10% overestimation when $X = 1$. All calibration parameters are learned on a balanced training distribution ($P(X = 0) = 0.5$). Prevalence is estimated by averaging the predicted probabilities over the target sample. We then evaluate prevalence estimates on test distributions where $P(X = 0)$ ranges from 0.01 to 0.99, repeating 50 times (details in Materials and Methods).

³An open-source implementation is available at <https://mcgrad.dev>.




Figure 1: Prevalence estimation bias (% relative) under covariate shift, averaged over 50 simulation runs. The x-axis shows $\Delta P(X=0)$, the change in $P(X=0)$ from the training value of 0.5. Classify & Count and Rogan-Gladen diverge with increasing shift (arrows indicate values beyond the axis range); isotonic regression (global calibration) shows moderate bias; MCGrad maintains near-zero bias across all shift levels. Additional methods (PACC, SLD, uncalibrated) shown in SI Appendix Figure S3.

Figure 1 shows the results. At the training distribution (center), all methods produce approximately unbiased estimates. As the distribution shifts, the methods diverge. Rogan-Gladen is particularly unstable: its ratio structure amplifies estimation errors, with bias exceeding $\pm 40\%$ at extreme shifts (arrows in Figure 1 indicate values beyond the axis range). Classify & Count shows growing bias in the same direction. Isotonic regression (global calibration) shows more moderate but still substantial bias (up to $\pm 15\%$). SLD and PACC exhibit comparably large failures (SI Appendix Figure S3). The multicalibrated estimator maintains near-zero bias and the lowest RMSE across the entire range of distribution shifts (SI Appendix Figure S2), confirming that the bias reduction does not come at the cost of increased variance.

Empirical application: employment prevalence under age distribution shift

Before applying multicalibration to LLM-generated scores, which present additional challenges such as coarse discretization and non-standard score distributions, we first demonstrate the core mechanism in a controlled setting with a standard ML classifier. We analyze data from the American Community Survey (ACS), a large-scale annual survey conducted by the U.S. Census Bureau. The concept to be measured is the rate of employment. The measurement device is a logistic regression model of binary employment status, with 16 sociodemographic features including age, education, marital status, disability status, and citizenship.

Setup. We train a logistic regression classifier on data from eight U.S. states (TX, MI, PA, OH, IL, GA, NC, VA) across 2016–2018, totaling approximately 1.5 million training observations. The remaining in-distribution data is split into a calibration set ($n \approx 644,000$) and a test set ($n \approx 920,000$). Six additional states (CA, NY, FL, WA, AZ, CO) are held out for out-of-distribution (OOD) evaluation. We fit two post-hoc calibration methods on the calibration set: isotonic regression (global calibration) and MCGrad, a multicalibration algorithm that enforces calibration conditional on both categorical and numerical features. We compare five prevalence estimation methods (Figure 2): Classify & Count with

a prevalence-matched threshold, Rogan-Gladen adjustment, importance-weighted estimation (IPW), isotonic regression, and MCGrad. Additional methods (PACC, SLD, raw scores) are reported in SI Appendix Table S1.

Age distribution shift. Employment rates vary dramatically by age: approximately 47% for ages 16–24, 76% for ages 25–54, 61% for ages 55–64, and 17% for ages 65+. This makes age an ideal dimension along which to construct meaningful distribution shifts. We create synthetic target populations by resampling test data with shifted age distributions: young-skewed (oversampling ages 16–30), old-skewed (oversampling ages 60+), and bimodal (oversampling both tails). The resulting populations have true employment rates ranging from 12.8% to 46.0%. All calibration parameters are estimated once on the original calibration set and held fixed across scenarios.




Figure 2: Prevalence estimation $|bias|$ (percentage points) under synthetic age distribution shift, for in-distribution states (left) and out-of-distribution states (right). Each marker shape represents a different age shift scenario; horizontal lines show the mean across scenarios. Full numerical results in SI Appendix Table S1.

Results. With no age shift, all methods produce approximately unbiased estimates (Figure 2). Under shift, the methods diverge sharply. Rogan-Gladen exhibits catastrophic failure (12–19pp bias under age shift). SLD, designed for label shift rather than covariate shift, shows comparably large bias (SI Appendix Table S1). Classify & Count and isotonic regression show moderate but growing bias (2–8pp). IPW performs well on simple shifts (≤ 1.2 pp) but fails on the bimodal shift (+4.7pp in-distribution, +6.3pp OOD) where the density ratio is hard to model.

MCGrad produces near-zero bias across all in-distribution scenarios (≤ 0.27 pp), including the bimodal shift where IPW struggles. In the OOD setting (held-out states with age shift), MCGrad’s bias increases modestly (0.88–1.35pp), reflecting geographic shift along an uncalibrated dimension. Even so, MCGrad maintains the lowest bias and RMSE across all scenarios.

LLM-based topic classification under cross-national shift

The ACS application uses a standard ML classifier with well-distributed probability estimates. We now examine the setting that motivates this paper: using an LLM as a zero-shot measurement device across multiple countries and languages. We compare two LLM output modes that reflect current practice and the confidence elicitation literature, respectively: (1) discrete Yes/No classifications, which is how LLMs are used in all of the applied studies cited in the introduction, and (2) probability scores obtained via direct probability elicitation, where the LLM estimates $P(\text{Yes})$ and $P(\text{No})$ without first committing to an answer. This design lets us test whether confidence elicitation improves prevalence estimation, and whether multicalibration is needed in either case.

Setup. We use Claude Opus 4.6 to classify 30,000 political texts from the Comparative Agendas Project (Baumgartner et al. 2006) as related to Law & Crime (CAP major topic code 12). The data comprise six sub-populations across four countries and four languages (5,000 documents each): Danish parliamentary questions, Spanish oral questions, U.S. congressional bills, Belgian newspaper articles, Spanish media articles, and Belgian TV news. Each document is classified twice in independent campaigns: once for a binary Yes/No label, and once for direct probability estimates $P(\text{Yes})$ and $P(\text{No})$ (summing to 1.0). The LLM achieves strong discriminative performance: AUC 0.959 for binary labels and 0.988 for probability scores on the in-distribution test set, with per-language AUCs of 0.978–0.985 for the probability scores.

The calibration set ($n \approx 13,400$) draws equally from four sub-populations (Denmark questions, Spain questions, U.S. bills, Belgium newspaper), ensuring that all four countries and languages are represented. Two sub-populations are held out as out-of-distribution targets: Spanish media and Belgian TV news, which share country and language with calibration data but introduce an unseen document type. MCGrad is calibrated with categorical features (country, document type, party) and two numerical features (decade, document length).

For the binary-label condition, MCGrad receives the LLM’s Yes/No classification as a categorical input feature, with all initial scores set to the calibration-set base rate. MCGrad then learns feature-conditional prevalence estimates from the label and metadata alone, without any probability score from the LLM. For the probability-score condition, MCGrad receives the LLM’s $P(\text{Yes})$ directly as the input score and calibrates it conditional on the same metadata features.

Results. Figure 3 shows prevalence estimation bias across five scenarios: a baseline with no shift, two within-calibration shifts (country composition, document type composition), and two OOD scenarios.




Figure 3: Prevalence estimation $|bias|$ (percentage points) across the shift gradient. Each marker shape represents a different shift scenario; horizontal lines show the mean across scenarios. MCGrad achieves near-zero bias within the calibration distribution in both binary-label and probability-score conditions. Full numerical results in SI Appendix Table S2.

The broken workflow is clearly visible: Classify & Count, the standard practice of counting positive LLM classifications, produces bias of +1.6 to +4.8pp across all scenarios. The Rogan-Gladen adjustment reduces bias within calibration but still shows +3.1 to +3.7pp on OOD populations. Isotonic regression shows moderate bias within calibration (≤ 1.1 pp) but larger OOD bias (2.4–2.7pp). SLD, designed for label shift, fails catastrophically (+5 to +14pp; SI Appendix Table S2).

MCGrad on binary labels achieves near-zero bias within the calibration distribution (≤ 0.5 pp) and degrades gracefully on OOD targets: +0.4pp on Belgian TV and -2.5pp on Spanish media. This is achieved using only the LLM’s Yes/No classification and document metadata, with no probability scores. MCGrad on probability scores achieves comparable within-calibration performance (≤ 0.2 pp) but shows larger

OOD bias on Spanish media (-4.4pp) and Belgian TV (+1.6pp). The finding that binary labels with MCGrad can match or outperform probability scores with MCGrad suggests that the metadata features, not the input score quality, drive the calibration improvement.

IPW, the standard covariate-shift method, achieves near-zero within-calibration bias but fails severely on OOD populations (-12.1pp on Spanish media) because the target contains feature values absent from the source (a positivity violation). Unlike MCGrad, IPW also requires re-estimating density ratios for each target population.

Comparison with the ACS application. Across both applications, the pattern is consistent: MCGrad achieves near-zero bias when the target population’s features are within the calibration support, and degrades when the shift is along an uncalibrated dimension. IPW matches MCGrad within calibration but fails harder on OOD targets. MCGrad’s practical advantage is that it requires no target-specific estimation: a single calibrated device can be applied to any target population. This is the “universal adaptability” property of Kim et al. (2022). We replicate the CAP analysis using Llama 3.3 70B Instruct (an open-weight model) in SI Appendix S2; the results are consistent, confirming that the findings are not specific to a particular LLM.

Discussion

We have shown that multicalibration, calibration conditional on the input features rather than just on average, is sufficient for accurate model-based prevalence estimation under population shift. The minimal theoretical requirement is multi-accuracy (correct predictions on average within each subgroup), which multicalibration implies. Both conditions require that the calibration features capture the relevant dimensions of population shift, and both require that the target population’s features lie within the support of the calibration distribution. Both empirical applications confirm this: when calibrated features overlap with the target population, MCGrad achieves near-zero bias; when the target introduces a novel feature value, bias increases with the severity of the shift. Practitioners should ensure their calibration data covers the key dimensions along which target populations may differ.

A central practical advantage of multicalibration over traditional covariate-shift methods is that it produces a *target-independent* measurement device. Importance-weighted methods achieve comparable within-calibration performance but require estimating density ratios for each new target population, and our CAP results show they fail more severely on OOD targets (IPW: -12.1pp on Spanish media, vs. MCGrad: -2.5pp) due to positivity violations when target features are outside the source support. A multicalibrated device, by contrast, is calibrated once and deployed without target-specific re-estimation. This is the “universal adaptability” property of Kim et al. (2022), and it is what makes multicalibration particularly well-suited to the measurement setting, where a single device is often applied to many target populations.

A notable finding from the CAP application is that MCGrad on discrete binary labels achieves comparable or better prevalence estimation than MCGrad on continuous probability scores. This suggests that for prevalence estimation under covariate shift, the metadata features that capture the shift structure matter more than the quality of the LLM’s confidence scores. This has immediate practical implications: researchers using the standard workflow of prompting an LLM for Yes/No classifications can apply multicalibration directly to those labels, without needing to extract probability scores via confidence elicitation methods.

The practical implications are immediate. Researchers using LLMs as measurement devices typically validate by reporting accuracy or AUC, which provide no information about calibration. A model with 95% accuracy can still produce heavily biased prevalence estimates if its errors are unevenly distributed across the feature space (Grimmer et al. 2022). Multicalibration does not require access to a model’s

internals: observable metadata (document source, language, text length) can serve directly as segment features (Detommaso et al. 2024), and as our CAP results show, even discrete Yes/No labels suffice as input.

Several limitations warrant discussion. First, the guarantee holds only for shifts along calibrated dimensions; OOD performance degrades when target populations introduce novel feature values. Second, multicalibration requires labeled calibration data, which may require the manual annotation LLMs were intended to avoid; prediction-powered inference (Angelopoulos et al. 2023) offers a complementary framework. Third, our framework assumes covariate shift ($P(Y | X)$ stable); under concept drift, no purely statistical correction substitutes for new labeled data. Finally, both applications use settings where ground-truth labels enable direct bias measurement.

These results connect two literatures that have developed in isolation. The quantification literature has focused on correcting aggregate error rates but has not engaged with feature-conditional calibration (González et al. 2017; Wu and Resnick 2024). The multicalibration literature has focused on individual-level prediction quality and fairness but has not emphasized the implications for population-level inference (Hébert-Johnson et al. 2018; Kim et al. 2022). Our contribution is to show that calibration conditional on the feature space is the missing link that makes LLM-based prevalence estimation reliable under the distribution shifts that motivate its use.

Materials and Methods

Simulation. We generate $n = 10,000$ observations with a binary covariate $X \in \{0, 1\}$ and outcome $Y \sim \text{Bernoulli}(P(Y | X))$, where $P(Y = 1 | X = 1) = 0.85$ and $P(Y = 1 | X = 0) = 0.15$. The classifier produces deterministic scores $\hat{p}(X = 0) = 0.135$ and $\hat{p}(X = 1) = 0.935$, representing 10% multiplicative bias within each stratum. For each of $B = 50$ iterations, we generate fresh calibration data from $P(X = 0) = 0.5$, estimate all method-specific parameters, then evaluate bias and mean squared error on 20 test distributions with $P(X = 0)$ ranging from 0.01 to 0.99. The seven methods compared (uncalibrated averaging, Classify & Count with a prevalence-matched threshold, Rogan-Gladen adjustment, PACC, SLD/EMQ, global multiplicative calibration, and multicalibration with stratum-specific additive corrections) are detailed with full mathematical definitions in the SI Appendix.

Empirical application (ACS). We use the ACS Public Use Microdata Sample via the *folktables* package, with the ACSEmployment prediction task (binary: employed vs. not employed) and 16 sociodemographic features. Training data comprises eight states (TX, MI, PA, OH, IL, GA, NC, VA) across 2016–2018. The base model is logistic regression with standard scaling. Post-hoc calibration uses isotonic regression (global) and MCGrad (multicalibration with categorical and numerical segment features) on a held-out calibration set ($n \approx 644,000$). In addition to Classify & Count, Rogan-Gladen, isotonic regression, and MCGrad, we include importance-weighted prevalence estimation (IPW), which estimates density ratios between source and target via logistic regression on all 16 features, and two methods from the quantification literature: PACC (Probabilistic Adjusted Classify & Count) (González et al. 2017) and SLD (Saerens-Latinne-Decaestecker) (Saerens et al. 2002). Raw (uncalibrated) scores are also reported in SI Appendix Table S1. Synthetic age-shifted populations are created by importance-weighted resampling of the test set ($n \approx 920,000$), with exponential weights favoring young ages, old ages, or both extremes. RMSE is computed via bootstrap resampling (200 iterations per scenario). Six additional states (CA, NY, FL, WA, AZ, CO) serve as OOD evaluation data.

Empirical application (CAP). We use data from the Comparative Agendas Project (Baumgartner et al. 2006), which provides expert-coded policy topic labels for political texts across countries. The binary outcome is whether a document addresses Law & Crime (CAP major topic code 12). We sample 30,000 documents (5,000 from each of six sub-populations: Danish parliamentary questions, Spanish oral

questions, U.S. congressional bills, Belgian newspaper articles, Spanish media articles, and Belgian TV news). The measurement device is Claude Opus 4.6. Each document is classified in two independent campaigns: (1) binary Yes/No classification using the CAP codebook definition of Law & Crime, and (2) direct probability elicitation, where the LLM estimates $P(\text{Yes})$ and $P(\text{No})$ without first committing to an answer. The two campaigns were run independently to avoid anchoring contamination. The calibration set ($n \approx 13,400$) draws equally from four sub-populations (Denmark questions, Spain questions, U.S. bills, and Belgium newspaper). MCGrad is calibrated with categorical features (country, document type, political party) and two numerical features (decade and document length). For the binary-label condition, the LLM’s Yes/No classification is passed to MCGrad as a categorical input feature, with initial scores set to the calibration-set base rate; MCGrad learns feature-conditional prevalence estimates from the metadata and LLM label alone. For the probability-score condition, MCGrad receives $P(\text{Yes})$ as the input score. Classify & Count reports the fraction of positive LLM classifications. Rogan-Gladen adjustment corrects this fraction using binary-label TPR and FPR estimated on the calibration set. Isotonic regression is fitted on probability scores using the same calibration set. IPW estimates density ratios via logistic regression on country, document type, decade, and document length. SLD results are reported in SI Appendix Table S2. Two sub-populations are held out as out-of-distribution targets: Spanish media and Belgian TV news, which differ from the calibration data only in document type while sharing the same countries and languages.

Software. MCGrad is available at github.com/facebookincubator/MCGrad. Simulation and analysis code are available at [TODO: insert repository URL before submission].

Author Contributions

[TODO: Add author contributions, e.g., “A.B. designed research; C.D. performed research; E.F. contributed analytic tools; A.B. and C.D. analyzed data; and A.B. wrote the paper.”]

Competing Interests

The authors declare no competing interests.

Data Availability

The American Community Survey data are publicly available via the *folkttables* package. The Comparative Agendas Project data are publicly available at <https://www.comparativeagendas.net>. Simulation and analysis code are available at [TODO: insert repository URL].

Keywords

multicalibration, prevalence estimation, covariate shift, large language models, quantification

References

- Angelopoulos, Anastasios N., Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. 2023. “Prediction-Powered Inference.” *Proceedings of the National Academy of Sciences* 120 (33).
- Baumgartner, Frank R., Christoffer Green-Pedersen, and Bryan D. Jones. 2006. “Comparative Studies of Policy Agendas.” *Journal of European Public Policy* 13 (7): 959–74.

- Benoit, Kenneth, Scott De Marchi, Conor Laver, Michael Laver, and Jinshuai Ma. 2026. “Using Large Language Models to Analyze Political Texts Through Natural Language Understanding.” *American Journal of Political Science*, ahead of print. <https://doi.org/10.1111/ajps.70050>.
- Detommaso, Gianluca, Martin Bertran, Riccardo Fogliato, and Aaron Roth. 2024. “Multicalibration for Confidence Scoring in LLMs.” *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Gilardi, Fabrizio, Meysam Alizadeh, and Maël Kubli. 2023. “ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks.” *Proceedings of the National Academy of Sciences* 120 (30): e2305016120.
- González, Pablo, Alberto Castaño, Nitesh V. Chawla, and Juan José Del Coz. 2017. “A Review on Quantification Learning.” *ACM Computing Surveys* 50 (5): 1–40.
- Grimmer, Justin, Margaret E. Roberts, and Brandon M. Stewart. 2022. “Text as Data: A New Framework for Machine Learning and the Social Sciences.” *Annual Review of Political Science* 24: 391–415.
- Hébert-Johnson, Ursula, Michael P. Kim, Omer Reingold, and Guy N. Rothblum. 2018. “Multicalibration: Calibration for the (Computationally-Identifiable) Masses.” *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 1939–48.
- Kadavath, Saurav, Tom Conerly, Amanda Askell, et al. 2022. “Language Models (Mostly) Know What They Know.” *arXiv Preprint arXiv:2207.05221*.
- Karjus, Andres et al. 2025. “Mixed-Methods Quantitizing with Prompt-Coded Variables Using LLMs.” *Humanities & Social Sciences Communications*.
- Kim, Michael P., Christoph Kern, Shafi Goldwasser, Frauke Kreuter, and Omer Reingold. 2022. “Universal Adaptability: Target-Independent Inference That Competes with Propensity Scoring.” *Proceedings of the National Academy of Sciences* 119 (4): e2108097119.
- Maletzke, André G., Denis M. dos Reis, Everton A. Cherman, and Gustavo E. A. P. A. Batista. 2020. “On the Need for Class Ratio Insensitive Drift Tests in Data Stream Mining.” *Proceedings of the 2nd International Workshop on Learning with Imbalanced Domains: Theory and Applications (LIDTA)*.
- Mellon, Jonathan et al. 2024. “Coding Open-Ended Survey Responses with Language Models at Scale.” *Research & Politics*.
- Oliveira, Rodrigo de, Matthew Garber, James M. Gwinnutt, et al. 2025. “A Study of Calibration as a Measurement of Trustworthiness of Large Language Models in Biomedical Natural Language Processing.” *JAMIA Open* 8 (4).
- Overös, Harris et al. 2024. “Zero-Shot Protest Event Detection with GPT Models.” *PNAS Nexus*.
- Ren, Kevin, Santiago Cortes-Gomez, Carlos Miguel Patino, et al. 2025. “Predicting Language Models’ Success at Zero-Shot Probabilistic Prediction.” *Findings of the Association for Computational Linguistics: EMNLP 2025*, 18337–63.
- Rogan, Walter J., and Beth Gladen. 1978. “Estimating Prevalence from the Results of a Screening Test.” *American Journal of Epidemiology* 107 (1): 71–76.

- Saerens, Marco, Patrice Latinne, and Christine Decaestecker. 2002. “Adjusting the Outputs of a Classifier to New a Priori Probabilities: A Simple Procedure.” *Neural Computation* 14 (1): 21–41.
- Sibley, James et al. 2025. “Assessment of a Zero-Shot Large Language Model in Measuring Documented Goals-of-Care Discussions.” *medRxiv Preprint*, ahead of print. <https://doi.org/10.1101/2025.05.23.25328115>.
- Storkey, Amos. 2009. “When Training and Test Sets Are Different: Characterizing Learning Transfer.” In *Dataset Shift in Machine Learning*. MIT Press.
- Sushil, Madhumita et al. 2024. “Zero-Shot Clinical Concept Extraction from Pathology Reports Using GPT-4.” *Journal of the American Medical Informatics Association (JAMIA)*.
- Tax, Niek, Lorenzo Perini, Fridolin Linder, et al. 2026. “MCGrad: Multicalibration at Web Scale.” *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*.
- Tian, Katherine, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, and Chelsea Finn. 2023. “Just Ask for Calibration: Strategies for Eliciting Calibrated Confidence Scores from Language Models Fine-Tuned with Human Feedback.” *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Tojima, Tatsuya, and Mitsuo Yoshida. 2025. “Zero-Shot Classification of Art with Large Language Models.” *IEEE Access* 13.
- Wang, Xuezhi, Jason Wei, Dale Schuurmans, et al. 2023. “Self-Consistency Improves Chain of Thought Reasoning in Language Models.” *Proceedings of the 11th International Conference on Learning Representations (ICLR)*.
- Weidmann, Nils B., Valentin Faulborn, and Estefanía Garcia. 2025. “Democracy Indicators Coded Zero-Shot with LLMs.” *PS: Political Science & Politics*.
- Wu, Lingjiao, and Paul Resnick. 2024. “Calibrate-Extrapolate: Rethinking Prevalence Estimation with Black Box Classifiers.” *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM)* 18.
- Yang, Yahan, Soham Dan, Dan Roth, and Insup Lee. 2023. “On the Calibration of Multilingual Question Answering LLMs.” *arXiv Preprint arXiv:2311.07009*.
- Yuan, Weizhe, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. 2024. “Self-Rewarding Language Models.” *arXiv Preprint arXiv:2401.10020*.
- Zheng, Lianmin, Wei-Lin Chiang, Ying Sheng, et al. 2023. “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena.” *Advances in Neural Information Processing Systems (NeurIPS)*.